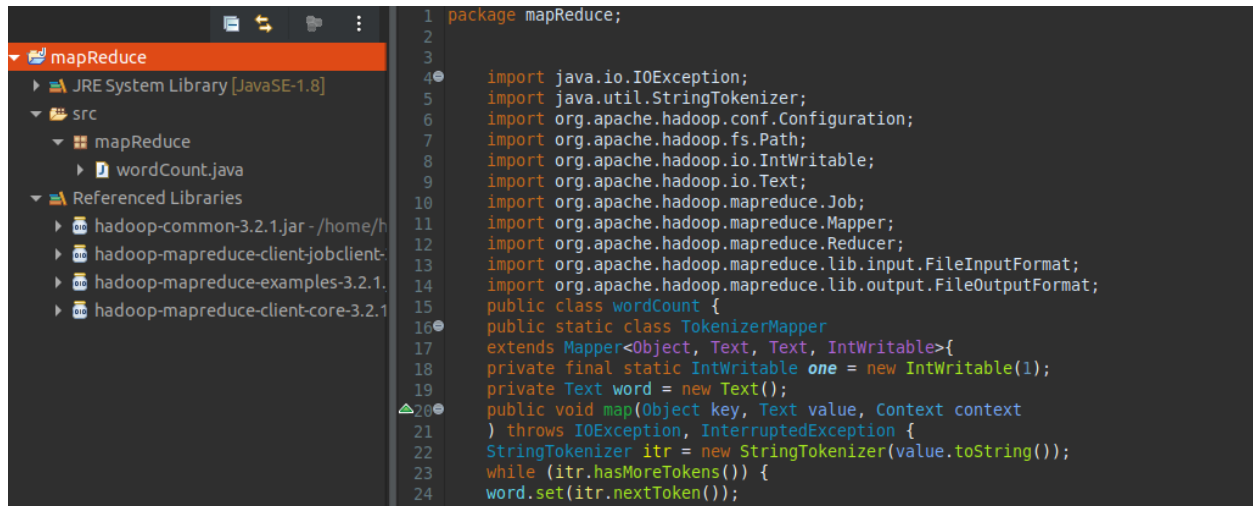


# Map Reduce

NAME: T H LAKSHANA

USN: 1NT19IS174



```
1 package mapReduce;
2
3
4 import java.io.IOException;
5 import java.util.StringTokenizer;
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 public class wordCount {
16     public static class TokenizerMapper
17     extends Mapper<Object, Text, Text, IntWritable>{
18         private final static IntWritable one = new IntWritable(1);
19         private Text word = new Text();
20     public void map(Object key, Text value, Context context
21         ) throws IOException, InterruptedException {
22         StringTokenizer itr = new StringTokenizer(value.toString());
23         while (itr.hasMoreTokens()) {
24             word.set(itr.nextToken());
25             context.write(word, one);
26         }
27     }
28 }
```

```
int sum = 0;
for (IntWritable val : values) {
    sum += val.get();
}
result.set(sum);
context.write(key, result);
}
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(wordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

```

hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1$ cd sbin
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ ls
1nt19is097          niki1nt19is097.txt      stop-all.sh
apache-hive-3.1.2-bin.tar.gz  p2.txt                  stop-balancer.sh
ch.txt              p3.txt                  stop-dfs.cmd
day.txt             refresh-namenodes.sh    stop-dfs.sh
distribute-exclude.sh  start-all.cmd          stop-secure-dns.sh
FederationStateStore  start-all.sh           stop-yarn.cmd
hadoop-daemon.sh      start-balancer.sh       stop-yarn.sh
hadoop-daemons.sh    start-dfs.cmd           workers.sh
httpfs.sh            start-dfs.sh            x.txt
kms.sh               start-secure-dns.sh     yarn-daemon.sh
mongodb-compass_1.26.1_amd64.deb  start-yarn.cmd          yarn-daemons.sh
mr-jobhistory-daemon.sh  start-yarn.sh
mytext.txt           stop-all.cmd
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ jps
5874 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
6506 Jps

```

```

hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [admin1-HP-280-G4-MT-Business-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -mkdir -p ~/input
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile ~/input/test.txt
appendToFile: stdin (-) must be the sole input argument when present
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile ~/input/test.txt
appendToFile: /home/hadoop (Is a directory)
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile ~/input/test.txt
my name is lakshana my name is greeshma2022-06-17 09:26:55,650 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted
= false, remoteHostTrusted = false

```

```

at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Desktop/lakshana.jar ~/input ~/out
2022-06-17 09:31:55,017 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-17 09:31:55,304 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and
execute your application with ToolRunner to remedy this.
2022-06-17 09:31:55,343 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_165
5437936096_0001
2022-06-17 09:31:55,457 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-17 09:31:55,556 INFO input.FileInputFormat: Total input files to process : 1
2022-06-17 09:31:55,636 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-17 09:31:55,696 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-17 09:31:55,718 INFO mapreduce.JobSubmitter: number of splits:1
2022-06-17 09:31:55,826 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-17 09:31:56,252 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1655437936096_0001
2022-06-17 09:31:56,252 INFO mapreduce.JobSubmitter: Executing with tokens: []

```

```

hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -cat ~/out/part*
2022-06-17 09:32:46,136 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
greeshma      1
is            2
lakshana      1
my            2
name          2

```